

# ***Software Archery: hitting the bull's-eye with GNU Arch***

**Cameron Patrick**

Perth Linux Users' Group  
March 2005



# *Free software revision control*

- Keeping your code organised
- Allowing easy access to past versions
- Coordinating changes between developers
- Allows working in different 'branches'
  - **and merging between them**
- So many revision control systems to choose from these days.



# CVS

- Oldest and most popular revision control system for free software.
  - Technical problems: doesn't track renames, symbolic links, file permissions, has issues with binary files.
  - Social problems: only core developers can use CVS, everyone else is left on their own. Branching and merging is tricky.
- 
-

# *Simple revision control*

- `cp -r myproject myproject.backup`
  - Every time you make a major change to your software, take a backup copy and name it consistently: e.g. v1, v2, v3, v4, ...
  - This is the mental model used by Subversion: every time you commit, the repository-wide version number increases.
  - Fine for a centralised system, but what if you want to send your changes to someone else?
- 
-

# *Simple revision control II*

- `diff -urN myproject.orig myproject`
- Keep a bunch of diffs lying around for each interesting change you make, and you can build the latest version from an older version and some patches.
- This is the mental model used by Arch.

# *Changesets*

- Diff and patch have some limitations: e.g. they don't handle file renames or binary files.
- Arch uses a 'changeset' instead: a bunch of patch files and some meta-data to handle things that patch can't, all wrapped up in a tar.gz file.

# ***Renames and logical file IDs***

- When a file is added to an Arch project, it's given a logical file identity.
- Inside a changeset we record a list of file names and their corresponding IDs.
- This means that if you have a file 'foo.c' but I've renamed it to 'silly\_foo.c', we can still exchange Arch changesets.



# *Turning it into a revision control system*

- Tar up the initial state of your project and keep it in a directory called `base-0`.
- Put changesets in numbered directories: `patch-1`, `patch-2`, `patch-3`, etc.
- We're now most of the way to having an Arch archive.





## *Aside: signed archives*

- Every changeset is stored in a separate file and never changes once it's been written.
- So it's trivial to store a GPG signature next to each changeset. (A *signed* Arch archive.)
- This means that if your server is compromised, you will be able to easily see if your code has been tampered with.
  - c.f. break-in at freedesktop.org not long ago

# *Distributed revision control: the Arch namespace*

- Give each changeset a globally unique name
  - e.g.: cp@chem.com.au--2005/  
hibernate--debian--0--patch-5
- Keep track of which changesets have already been applied in each tree
- Now we can automatically merge changes made by someone else, in a completely different Arch archive
  - apply the changesets that we don't already have

# *Arch development: history and future*

- Arch was original implemented as a shell script known as `larch`. Don't use it!
  - Rewritten in C as `tla` (Tom Lord's Arch)
    - `tla 1.x` is the production version
    - `tla 2.0` is an ambitious rewrite, just begun
  - Forked by Canonical (the folks behind Ubuntu) as `bazaar` or `baz`
    - Making more aggressive user interface changes and testing some of the new ideas for `tla 2.0`
- 
-

# *Some reasons why clever merging is cool*

- Hacking on someone else's project
  - making your own changes in your own archive
  - catching up with their changes as necessary
  - and they can easily merge some or all of your changes
- Working off-line on a laptop
  - committing to an archive stored on your laptop
  - push changes back to your main archive when you're back at home



# *Sharing your code with the world*

- No special “arch server” is needed.
    - Arch doesn’t even need to be installed on the server machine.
    - Your choice of SSH/SFTP, HTTP / WebDAV, or straight ftp.
  - Usual combination is SFTP for committing and HTTP for read access
  - ‘tla archive-mirror’ command to copy new revisions from your archive to a remote server
- 
-

# *Getting started with arch*

- Download Bazaar from <http://bazaar.canonical.com/>
  - Debian/Ubuntu packages are available
- Then head over to [wiki.gnuarch.org](http://wiki.gnuarch.org)
  - read the “quick introduction” and/or “learning Arch for CVS users”
- Maybe join the gnu-arch-users mailing list



# ***Demonstration***



# *"Arch sucks, what else is there?"*

- Lots of free distributed revision control systems to choose from.
    - Arch is the most popular, but that doesn't necessarily mean it's the best for you.
  - Darcs
  - Monotone
  - Quilt (not really a revision control system)
  - Subversion (not really distributed)
  - Bazaar-NG (still in very early development)
- 
-



# ***Darcs (David's Advanced Revision Control System)***

- Amazingly simple to learn and become productive with
    - Though some operations which are simple in arch and cvs are harder in Darcs
  - Scalability issues with
    - large projects
    - lots of branches
  - One to watch: rapidly gaining momentum and catching up to Arch
- 
-

# *Questions?*

